

CS539 Assignment 04

Part 2

When the Network only has two fully connected network, that is **without CNN** layer, we got the results as follows:

```
[107]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = Net().to(device)

optimizer = optim.SGD(model.parameters(), lr=args['lr'], momentum=args['momentum'])

for epoch in range(1, args['epochs'] + 1):
    train(epoch)
    test()
```

Train Epoch: 1	[56640/60000 (94%)]	Loss: 0.285482
Train Epoch: 1	[56960/60000 (95%)]	Loss: 0.054643
Train Epoch: 1	[57280/60000 (95%)]	Loss: 0.015016
Train Epoch: 1	[57600/60000 (96%)]	Loss: 0.023588
Train Epoch: 1	[57920/60000 (97%)]	Loss: 0.188009
Train Epoch: 1	[58240/60000 (97%)]	Loss: 0.290808
Train Epoch: 1	[58560/60000 (98%)]	Loss: 0.135150
Train Epoch: 1	[58880/60000 (98%)]	Loss: 0.015152
Train Epoch: 1	[59200/60000 (99%)]	Loss: 0.104854
Train Epoch: 1	[59520/60000 (99%)]	Loss: 0.047242
Train Epoch: 1	[59840/60000 (100%)]	Loss: 0.048344

```
/opt/anaconda3/lib/python3.12/site-packages/torch/nn/_reduction.py:42: UserWarning: size_average and reduce args will be deprecated, please
se reduction='sum' instead.
  warnings.warn(warning.format(ret))

Test set: Average loss: 0.0990, Accuracy: 9693/10000 (97%)
```

Average Loss: 0.0990

Accuracy: 9693/10000 (97%)

When the network has 1 CNN layer with pooling and two FC layers, that is **with CNN layer**, it resulted as:

```
[219]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = NetWithCNN().to(device)

optimizer = optim.SGD(model.parameters(), lr=args['lr'], momentum=args['momentum'])

#optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

for epoch in range(1, args['epochs'] + 1):
    train(epoch)
    test()
```

Train Epoch: 1	[53120/60000 (88%)]	Loss: 0.150327
Train Epoch: 1	[53760/60000 (90%)]	Loss: 0.159112
Train Epoch: 1	[54400/60000 (91%)]	Loss: 0.186541
Train Epoch: 1	[55040/60000 (92%)]	Loss: 0.187486
Train Epoch: 1	[55680/60000 (93%)]	Loss: 0.214006
Train Epoch: 1	[56320/60000 (94%)]	Loss: 0.270832
Train Epoch: 1	[56960/60000 (95%)]	Loss: 0.124888
Train Epoch: 1	[57600/60000 (96%)]	Loss: 0.097894
Train Epoch: 1	[58240/60000 (97%)]	Loss: 0.109332
Train Epoch: 1	[58880/60000 (98%)]	Loss: 0.082911
Train Epoch: 1	[59520/60000 (99%)]	Loss: 0.152225

```
/opt/anaconda3/lib/python3.12/site-packages/torch/nn/_reduction.py:42: UserWarning: size_average and reduce args will be deprecated, please
se reduction='sum' instead.
  warnings.warn(warning.format(ret))

Test set: Average loss: 0.1585, Accuracy: 9554/10000 (96%)
```

Average Loss: 0.1585

Accuracy: 9554/10000 (96%)

Experimenting with 3 alternative network topologies and hyper-parameters as:

#1. 2 CNN layers + 2 FC layers

I used 2 CNN layers with pooling and 2 fully – connected (FC) layers.

```
[231]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = NetExp1().to(device)

optimizer = optim.SGD(model.parameters(), lr=args['lr'], momentum=args['momentum'])
#optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

for epoch in range(1, args['epochs'] + 1):
    train(epoch)
    test()
```

Train Epoch: 1	[50560/60000 (84%)]	Loss: 0.060764
Train Epoch: 1	[51200/60000 (85%)]	Loss: 0.171581
Train Epoch: 1	[51840/60000 (86%)]	Loss: 0.300890
Train Epoch: 1	[52480/60000 (87%)]	Loss: 0.142780
Train Epoch: 1	[53120/60000 (88%)]	Loss: 0.265829
Train Epoch: 1	[53760/60000 (90%)]	Loss: 0.155733
Train Epoch: 1	[54400/60000 (91%)]	Loss: 0.223365
Train Epoch: 1	[55040/60000 (92%)]	Loss: 0.053974
Train Epoch: 1	[55680/60000 (93%)]	Loss: 0.130385
Train Epoch: 1	[56320/60000 (94%)]	Loss: 0.217469
Train Epoch: 1	[56960/60000 (95%)]	Loss: 0.138121
Train Epoch: 1	[57600/60000 (96%)]	Loss: 0.213913
Train Epoch: 1	[58240/60000 (97%)]	Loss: 0.137802
Train Epoch: 1	[58880/60000 (98%)]	Loss: 0.168160
Train Epoch: 1	[59520/60000 (99%)]	Loss: 0.169995

Test set: Average loss: 0.1690, Accuracy: 9494/10000 (95%)

Average Loss: 0.1690

Accuracy: 9494/10000 (95%)

#2. Learning Rate Tuning with 1 CNN layer model

I tested with 3 different learning rates with a simplified model (1 CNN layer):

Learning Rate	Test Loss (Avg Loss)	Test Accuracy (in %)	Notes
0.01	0.0585	97.90%	Best balance of speed/accuracy
0.001	0.1351	95.34%	Too slow convergence
0.5	2.3293	10.28%	Diverged

Best LR = 0.01: Achieved 97.9 % accuracy, the highest among the tested LR.

#3. Batch Size Tuning

I tested with 3 different batch sizes with the model having 1 CNN layer with pooling and 2 FC layers with LR = 0.01 and momentum = 0.9

Batch Size	Test Loss (avg. loss)	Test Accuracy	Notes
16	0.0532	98.31%	Highest accuracy
32	0.0581	98.03%	Default; stable
64	0.0723	97.63%	Slightly higher loss

The best configuration was:

1 CNN + 2 FC with the Test accuracy as 98.31% having LR = 0.01 and batch size = 16

Through me experimenting, I learned that **hyperparameter tuning significantly impacts model performance**. First, the **learning rate** must be carefully chosen—too high (0.5) causes divergence, while too low (0.001) slows convergence, with 0.01 proving optimal for stable training. Second, **batch size** affects both accuracy and training dynamics: smaller batches (16) yielded the highest accuracy (98.31%) due to noisier but more exploratory updates, while larger batches (64) traded slight accuracy loss (97.63%) for faster training. Finally, **model complexity** isn't always better—a simpler 1-CNN-layer model outperformed a deeper 2-CNN-layer one (98.31% vs. 96.12%), suggesting that overly complex architectures can overfit even on simple datasets like MNIST. These results highlight the importance of **balanced hyperparameters, trade-offs between speed and precision**, and the **risk of over-engineering** models.